

```
#Import necessary libraries and modules from TensorFlow and Keras
```

```
import numpy as np
from tensorflow import keras
from tensorflow.keras.datasets import imdb
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Embedding, SimpleRNN, Dense, LSTM
from tensorflow.keras.preprocessing import sequence
```

```
#Download and extract the IMDB dataset.
```

```
!curl -O https://ai.stanford.edu/~amaas/data/sentiment/aclImdb_v1.tar.gz
!tar -xf aclImdb_v1.tar.gz
```

```

% Total    % Received % Xferd  Average Speed   Time    Time     Time  Current
           Dload  Upload   Total   Spent    Left     Speed
100 80.2M  100 80.2M    0     0  9244k      0  0:00:08  0:00:08 --:--:-- 16.8M

```

```
#Define the maximum number of features and the maximum length of sequences
```

```
max_features = 10000
maxlen = 150
(x_train, y_train), (x_test, y_test) = imdb.load_data(num_words=max_features)
x_train = sequence.pad_sequences(x_train, maxlen=maxlen)
x_test = sequence.pad_sequences(x_test, maxlen=maxlen)
```

```

Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/imdb.npz
17464789/17464789 ————— 0s 0us/step

```

```
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import layers
```

```
#Custom layer representing one layer of a transformer encoder
```

```
class TransformerEncoder(layers.Layer):
    def __init__(self, embed_dim, dense_dim, num_heads, **kwargs):
        super().__init__(**kwargs)
        self.embed_dim = embed_dim
        self.dense_dim = dense_dim
        self.num_heads = num_heads
        self.attention = layers.MultiHeadAttention(
            num_heads=num_heads, key_dim=embed_dim)
```

```
#Dense projection network
```

```
self.dense_proj = keras.Sequential(
    [layers.Dense(dense_dim, activation="relu"),
     layers.Dense(embed_dim),]
)
self.layernorm_1 = layers.LayerNormalization()
self.layernorm_2 = layers.LayerNormalization()
```

```
#Define how inputs are processed by the layer
```

```
def call(self, inputs, mask=None):
    if mask is not None:
        mask = mask[:, tf.newaxis, :]
    attention_output = self.attention(
        inputs, inputs, attention_mask=mask)

    proj_input = self.layernorm_1(inputs + attention_output)
    proj_output = self.dense_proj(proj_input)
    return self.layernorm_2(proj_input + proj_output)
```

```
#Serialize the layer's configuration
```

```
def get_config(self):
    config = super().get_config()
    config.update({
        "embed_dim": self.embed_dim,
        "num_heads": self.num_heads,
        "dense_dim": self.dense_dim,
    })
    return config
```

```
#Define parameters for the Transformer encoder layer
```

```

vocab_size = 20000
embed_dim = 256
num_heads = 2
dense_dim = 32

```

```

#Defined a model architecture using an embedding layer followed by a Transformer encoder layer and other dense layers.
#Compiling the model with optimizer, loss function, and metrics

```

```

inputs = keras.Input(shape=(None,), dtype="int64")
x = layers.Embedding(vocab_size, embed_dim)(inputs)
x = TransformerEncoder(embed_dim, dense_dim, num_heads)(x)
x = layers.GlobalMaxPooling1D()(x)
x = layers.Dropout(0.5)(x)
outputs = layers.Dense(1, activation="sigmoid")(x)
model = keras.Model(inputs, outputs)
model.compile(optimizer="rmsprop",
              loss="binary_crossentropy",
              metrics=["accuracy"])
model.summary()

```

Model: "functional\_1"

Layer (type)	Output Shape	Param #
input_layer (InputLayer)	(None, None)	0
embedding (Embedding)	(None, None, 256)	5,120,000
transformer_encoder (TransformerEncoder)	(None, None, 256)	543,776
global_max_pooling1d (GlobalMaxPooling1D)	(None, 256)	0
dropout_1 (Dropout)	(None, 256)	0
dense_2 (Dense)	(None, 1)	257

Total params: 5,664,033 (21.61 MB)

```

# Visualize sequence lengths before padding
import matplotlib.pyplot as plt

```

```

# Visualize sequence lengths before padding
sequence_lengths = [len(seq) for seq in imdb.load_data(num_words=max_features)[0][0]]


```

```

plt.hist(sequence_lengths, bins=30, color='blue', alpha=0.7)
plt.title("Distribution of Sequence Lengths Before Padding")
plt.xlabel("Sequence Length")
plt.ylabel("Frequency")
plt.show()

```

Distribution of Sequence Lengths Before Padding



```
import string

#Defining a class Vectorizer with a method "standardize" to preprocess text by converting it to lowercase and removing punctuation.
class Vectorizer:
    def standardize(self, text):
        text = text.lower()
        return "".join(char for char in text if char not in string.punctuation)
```

```
#Define a method "tokenize"
```

```
def tokenize(self, text):
    text = self.standardize(text)
    return text.split()
```

```
#creating a vocabulary from the dataset through a method "make_vocabulary"
```

```
def make_vocabulary(self, dataset):
    self.vocabulary = {"": 0, "[UNK]": 1}
    for text in dataset:
        text = self.standardize(text)
        tokens = self.tokenize(text)
        for token in tokens:
            if token not in self.vocabulary:
                self.vocabulary[token] = len(self.vocabulary)
    self.inverse_vocabulary = dict(
        (v, k) for k, v in self.vocabulary.items())
```

```
#Defining a method "encode" to encode text using the created vocabulary
```

```
def encode(self, text):
    text = self.standardize(text)
    tokens = self.tokenize(text)
    return [self.vocabulary.get(token, 1) for token in tokens]
```

```
#Decoding integer sequences back to text
```

```
def decode(self, int_sequence):
    return " ".join(
        self.inverse_vocabulary.get(i, "[UNK]") for i in int_sequence)
```

```
import re
import string
import tensorflow as tf
```

```
#Using a custom standardization function for text preprocessing using TensorFlow
```

```
def custom_standardization_fn(string_tensor):
    lowercase_string = tf.strings.lower(string_tensor)
    return tf.strings.regex_replace(
        lowercase_string, f"[{re.escape(string.punctuation)}]", "")
```

```
#Using a custom split function for text preprocessing
```

```
def custom_split_fn(string_tensor):
    return tf.strings.split(string_tensor)
```

```
#Removing the unsupervised training data directory, displaying the content of a positive review file
```

```
!curl -O https://ai.stanford.edu/~amaas/data/sentiment/aclImdb_v1.tar.gz
!tar -xf aclImdb_v1.tar.gz
!rm -r aclImdb/train/unsup
!cat aclImdb/train/pos/4077_10.txt
import os, pathlib, shutil, random
```

```
↻ % Total    % Received % Xferd  Average Speed   Time    Time     Time  Current
   100  80.2M  100  80.2M    0     0    9208k    0  0:00:08  0:00:08 --:--:-- 16.9M
I first saw this back in the early 90s on UK TV, i did like it then but i missed the chance to tape it, many years passed but the film a
```

```
#Defining batch size and directories for the dataset
```

```
batch_size = 32
base_dir = pathlib.Path("aclImdb")
val_dir = base_dir / "val"
train_dir = base_dir / "train"
```

```
#Moving a portion of training data to the validation directory for each category
#Here, I am making sure the directory and file exist
```

```
for category in ("neg", "pos"):
```

```
    if not os.path.exists(val_dir / category):
        os.makedirs(val_dir / category)
    files = os.listdir(train_dir / category)
    random.Random(1337).shuffle(files)
    num_val_samples = int(0.2 * len(files))
    val_files = files[-num_val_samples:]
    for fname in val_files:

        if not os.path.exists(val_dir / category / fname):
            shutil.move(train_dir / category / fname,
                        val_dir / category / fname)
```

```
from tensorflow import keras
batch_size = 32
```

```
#Creating TensorFlow datasets for training, validation, and testing
```

```
train_ds = keras.utils.text_dataset_from_directory(
    "aclImdb/train", batch_size=batch_size
)
val_ds = keras.utils.text_dataset_from_directory(
    "aclImdb/val", batch_size=batch_size
)
test_ds = keras.utils.text_dataset_from_directory(
    "aclImdb/test", batch_size=batch_size
)
```

```
→ Found 20000 files belonging to 2 classes.
   Found 5000 files belonging to 2 classes.
   Found 25000 files belonging to 2 classes.
```

```
#Displaying the shapes and dtypes of the first batch
```

```
for inputs, targets in train_ds:
```

```
    print("inputs.shape:", inputs.shape)
    print("inputs.dtype:", inputs.dtype)
    print("targets.shape:", targets.shape)
    print("targets.dtype:", targets.dtype)
    print("inputs[0]:", inputs[0])
    print("targets[0]:", targets[0])
    break
```

```
→ inputs.shape: (32,)
   inputs.dtype: <dtype: 'string'>
   targets.shape: (32,)
   targets.dtype: <dtype: 'int32'>
   inputs[0]: tf.Tensor(b'"Freddy\'s Dead" did the smartest thing it could\'ve done after the disappointment of the fifth film. It started
   targets[0]: tf.Tensor(0, shape=(), dtype=int32)
```

```
from tensorflow.keras.layers import TextVectorization
```

```
#Creating a dataset containing only text inputs for training
```

```
text_vectorization = TextVectorization(
    max_tokens=20000,
    output_mode="multi_hot",
)
text_only_train_ds = train_ds.map(lambda x, y: x)
text_vectorization.adapt(text_only_train_ds)
```

#Using the text\_vectorization function to convert text data into binary 1-gram representations

#Here, the use of "num\_parallel\_calls=4" argument indicates mapping done in parallel using four threads

```
binary_1gram_train_ds = train_ds.map(
    lambda x, y: (text_vectorization(x), y),
    num_parallel_calls=4)
binary_1gram_val_ds = val_ds.map(
    lambda x, y: (text_vectorization(x), y),
    num_parallel_calls=4)
binary_1gram_test_ds = test_ds.map(
    lambda x, y: (text_vectorization(x), y),
    num_parallel_calls=4)
```

#Data Exploration in binary\_1gram\_train\_ds

```
for inputs, targets in binary_1gram_train_ds:
    print("inputs.shape:", inputs.shape)
    print("inputs.dtype:", inputs.dtype)
    print("targets.shape:", targets.shape)
    print("targets.dtype:", targets.dtype)
    print("inputs[0]:", inputs[0])
    print("targets[0]:", targets[0])
    break
```

```
inputs.shape: (32, 20000)
inputs.dtype: <dtype: 'int64'>
targets.shape: (32,)
targets.dtype: <dtype: 'int32'>
inputs[0]: tf.Tensor([0 1 1 ... 0 0 0], shape=(20000,), dtype=int64)
targets[0]: tf.Tensor(0, shape=(), dtype=int32)
```

```
from tensorflow import keras
from tensorflow.keras import layers
```

#Function to Build and Compile a Feedforward Neural Network Model

#max\_tokens (int) indicates maximum number of tokens in each input, while sequence.hidden\_dim (int) indicates dimensionality of the hidden 1

```
def get_model(max_tokens=20000, hidden_dim=16):
    inputs = keras.Input(shape=(max_tokens,))
    x = layers.Dense(hidden_dim, activation="relu")(inputs)
    x = layers.Dropout(0.5)(x)
    outputs = layers.Dense(1, activation="sigmoid")(x)
    model = keras.Model(inputs, outputs)
    model.compile(optimizer="rmsprop",
                  loss="binary_crossentropy",
                  metrics=["accuracy"])
    return model
```

#Creating and summarizing a simple dense model.

```
model = get_model()
model.summary()
callbacks = [
    keras.callbacks.ModelCheckpoint("binary_1gram.keras",
                                    save_best_only=True)
]
```

Model: "functional\_2"

Layer (type)	Output Shape	Param #
input_layer_2 (InputLayer)	(None, 20000)	0
dense_3 (Dense)	(None, 16)	320,016
dropout_2 (Dropout)	(None, 16)	0
dense_4 (Dense)	(None, 1)	17

Total params: 320,033 (1.22 MB)  
Trainable params: 320,033 (1.22 MB)

```
#Training the model using the integer encoded training and validation dataset
```

```
model.fit(binary_1gram_train_ds.cache(),
          validation_data=binary_1gram_val_ds.cache(),
          epochs=10,
          callbacks=callbacks)
```

```
#Loading the model and evaluating its performance on the test dataset
```

```
model = keras.models.load_model("binary_1gram.keras")
print(f"Test acc: {model.evaluate(binary_1gram_test_ds)[1]:.3f}")
```

```
Epoch 1/10
625/625 ————— 10s 14ms/step - accuracy: 0.7682 - loss: 0.4865 - val_accuracy: 0.8900 - val_loss: 0.2794
Epoch 2/10
625/625 ————— 4s 6ms/step - accuracy: 0.8897 - loss: 0.2873 - val_accuracy: 0.8930 - val_loss: 0.2818
Epoch 3/10
625/625 ————— 4s 6ms/step - accuracy: 0.9116 - loss: 0.2463 - val_accuracy: 0.8946 - val_loss: 0.2955
Epoch 4/10
625/625 ————— 5s 6ms/step - accuracy: 0.9147 - loss: 0.2313 - val_accuracy: 0.8934 - val_loss: 0.3122
Epoch 5/10
625/625 ————— 5s 6ms/step - accuracy: 0.9243 - loss: 0.2155 - val_accuracy: 0.8926 - val_loss: 0.3245
Epoch 6/10
625/625 ————— 4s 7ms/step - accuracy: 0.9302 - loss: 0.2118 - val_accuracy: 0.8884 - val_loss: 0.3497
Epoch 7/10
625/625 ————— 4s 6ms/step - accuracy: 0.9331 - loss: 0.2064 - val_accuracy: 0.8924 - val_loss: 0.3482
Epoch 8/10
625/625 ————— 4s 6ms/step - accuracy: 0.9367 - loss: 0.2020 - val_accuracy: 0.8910 - val_loss: 0.3690
Epoch 9/10
625/625 ————— 6s 7ms/step - accuracy: 0.9331 - loss: 0.2032 - val_accuracy: 0.8900 - val_loss: 0.3697
Epoch 10/10
625/625 ————— 4s 6ms/step - accuracy: 0.9340 - loss: 0.2049 - val_accuracy: 0.8888 - val_loss: 0.3771
782/782 ————— 4s 5ms/step - accuracy: 0.8859 - loss: 0.2899
Test acc: 0.886
```

```
#Configuring TextVectorization Parameters
```

```
text_vectorization = TextVectorization(
    ngrams=2,
    max_tokens=20000,
    output_mode="multi_hot",
)
```

```
#Training and testing the binary bigram model
```

```
text_vectorization.adapt(text_only_train_ds)
binary_2gram_train_ds = train_ds.map(
    lambda x, y: (text_vectorization(x), y),
    num_parallel_calls=4)
binary_2gram_val_ds = val_ds.map(
    lambda x, y: (text_vectorization(x), y),
    num_parallel_calls=4)
binary_2gram_test_ds = test_ds.map(
    lambda x, y: (text_vectorization(x), y),
    num_parallel_calls=4)
```

```
#Initializing Model and Setting Up Checkpoints
```

```
model = get_model()
model.summary()
callbacks = [
    keras.callbacks.ModelCheckpoint("binary_2gram.keras",
                                    save_best_only=True)
]
```

```
Model: "functional_3"
```

Layer (type)	Output Shape	Param #
input_layer_3 (InputLayer)	(None, 20000)	0
dense_5 (Dense)	(None, 16)	320,016
dropout_3 (Dropout)	(None, 16)	0
dense_6 (Dense)	(None, 1)	17

```
Total params: 320,033 (1.22 MB)
Trainable params: 320,033 (1.22 MB)
```

#Training the Model and Evaluating Test Accuracy

```
model.fit(binary_2gram_train_ds.cache(),
          validation_data=binary_2gram_val_ds.cache(),
          epochs=10,
          callbacks=callbacks)
model = keras.models.load_model("binary_2gram.keras")
print(f"Test acc: {model.evaluate(binary_2gram_test_ds)[1]:.3f}")
```

```
Epoch 1/10
625/625 ————— 10s 14ms/step - accuracy: 0.7779 - loss: 0.4684 - val_accuracy: 0.8976 - val_loss: 0.2642
Epoch 2/10
625/625 ————— 6s 7ms/step - accuracy: 0.9046 - loss: 0.2623 - val_accuracy: 0.8988 - val_loss: 0.2661
Epoch 3/10
625/625 ————— 4s 6ms/step - accuracy: 0.9259 - loss: 0.2098 - val_accuracy: 0.8974 - val_loss: 0.2878
Epoch 4/10
625/625 ————— 4s 6ms/step - accuracy: 0.9390 - loss: 0.1975 - val_accuracy: 0.8970 - val_loss: 0.3077
Epoch 5/10
625/625 ————— 6s 8ms/step - accuracy: 0.9445 - loss: 0.1841 - val_accuracy: 0.8946 - val_loss: 0.3186
Epoch 6/10
625/625 ————— 4s 6ms/step - accuracy: 0.9498 - loss: 0.1770 - val_accuracy: 0.8946 - val_loss: 0.3242
Epoch 7/10
625/625 ————— 4s 6ms/step - accuracy: 0.9538 - loss: 0.1738 - val_accuracy: 0.8940 - val_loss: 0.3412
Epoch 8/10
625/625 ————— 5s 7ms/step - accuracy: 0.9532 - loss: 0.1590 - val_accuracy: 0.8942 - val_loss: 0.3530
Epoch 9/10
625/625 ————— 4s 6ms/step - accuracy: 0.9569 - loss: 0.1657 - val_accuracy: 0.8914 - val_loss: 0.3619
Epoch 10/10
625/625 ————— 4s 6ms/step - accuracy: 0.9561 - loss: 0.1618 - val_accuracy: 0.8914 - val_loss: 0.3759
782/782 ————— 6s 7ms/step - accuracy: 0.8962 - loss: 0.2714
Test acc: 0.896
```

#Configuring the TextVectorization layer to return token counts

```
text_vectorization = TextVectorization(
    ngrams=2,
    max_tokens=20000,
    output_mode="count"
)
```

#Configuring TextVectorization to return TF-IDF-weighted outputs

```
text_vectorization = TextVectorization(
    ngrams=2,
    max_tokens=20000,
    output_mode="tf_idf",
)
```

#Training and testing the TF-IDF bigram model

```
text_vectorization.adapt(text_only_train_ds)

tfidf_2gram_train_ds = train_ds.map(
    lambda x, y: (text_vectorization(x), y),
    num_parallel_calls=4)
tfidf_2gram_val_ds = val_ds.map(
    lambda x, y: (text_vectorization(x), y),
    num_parallel_calls=4)
tfidf_2gram_test_ds = test_ds.map(
    lambda x, y: (text_vectorization(x), y),
    num_parallel_calls=4)
```

#Initializing the Model and Summary

```
model = get_model()
model.summary()
```

#Model Checkpoint Configuration

```
callbacks = [
    keras.callbacks.ModelCheckpoint("tfidf_2gram.keras",
                                    save_best_only=True)
]
```

Model: "functional\_4"

Layer (type)	Output Shape	Param #
input_layer_4 (InputLayer)	(None, 20000)	0
dense_7 (Dense)	(None, 16)	320,016
dropout_4 (Dropout)	(None, 16)	0
dense_8 (Dense)	(None, 1)	17

Total params: 320,033 (1.22 MB)  
Trainable params: 320,033 (1.22 MB)

#Model Training and Evaluation on Test Data

```
model.fit(tfidf_2gram_train_ds.cache(),
          validation_data=tfidf_2gram_val_ds.cache(),
          epochs=10,
          callbacks=callbacks)
model = keras.models.load_model("tfidf_2gram.keras")
print(f"Test acc: {model.evaluate(tfidf_2gram_test_ds)[1]:.3f}")
```

Epoch 1/10  
625/625 ————— 9s 12ms/step - accuracy: 0.7295 - loss: 0.5449 - val\_accuracy: 0.8704 - val\_loss: 0.3144  
Epoch 2/10  
625/625 ————— 5s 7ms/step - accuracy: 0.8814 - loss: 0.2999 - val\_accuracy: 0.8880 - val\_loss: 0.2847  
Epoch 3/10  
625/625 ————— 4s 6ms/step - accuracy: 0.9007 - loss: 0.2573 - val\_accuracy: 0.8864 - val\_loss: 0.3006  
Epoch 4/10  
625/625 ————— 5s 6ms/step - accuracy: 0.9107 - loss: 0.2274 - val\_accuracy: 0.8902 - val\_loss: 0.3000  
Epoch 5/10  
625/625 ————— 6s 6ms/step - accuracy: 0.9195 - loss: 0.2156 - val\_accuracy: 0.8940 - val\_loss: 0.2943  
Epoch 6/10  
625/625 ————— 3s 6ms/step - accuracy: 0.9244 - loss: 0.2041 - val\_accuracy: 0.8892 - val\_loss: 0.3156  
Epoch 7/10  
625/625 ————— 4s 6ms/step - accuracy: 0.9238 - loss: 0.2015 - val\_accuracy: 0.8938 - val\_loss: 0.3161  
Epoch 8/10  
625/625 ————— 4s 7ms/step - accuracy: 0.9286 - loss: 0.1916 - val\_accuracy: 0.8952 - val\_loss: 0.3385  
Epoch 9/10  
625/625 ————— 3s 6ms/step - accuracy: 0.9326 - loss: 0.1894 - val\_accuracy: 0.8760 - val\_loss: 0.4114  
Epoch 10/10  
625/625 ————— 4s 6ms/step - accuracy: 0.9348 - loss: 0.1907 - val\_accuracy: 0.8912 - val\_loss: 0.3499  
782/782 ————— 6s 7ms/step - accuracy: 0.8901 - loss: 0.2849  
Test acc: 0.890

#Model Inference and Prediction

```
inputs = keras.Input(shape=(1,), dtype="string")
processed_inputs = text_vectorization(inputs)
outputs = model(processed_inputs)
inference_model = keras.Model(inputs, outputs)
import tensorflow as tf
raw_text_data = tf.convert_to_tensor([
    ["That was an excellent movie, I loved it."],
])
predictions = inference_model(raw_text_data)
print(f"{float(predictions[0] * 100):.2f} percent positive")
```

90.10 percent positive

#Processing words as a sequence: The sequence model approach

#Downloading the data

```
!curl -O https://ai.stanford.edu/~amaas/data/sentiment/aclImdb_v1.tar.gz
!tar -xf aclImdb_v1.tar.gz
!rm -r aclImdb/train/unsup
```

% Total	% Received	% Xferd	Average Speed	Time	Time	Time	Current
			Dload Upload	Total	Spent	Left	Speed
100 80.2M	100 80.2M	0 0	9080k 0	0:00:09	0:00:09	--:--:--	15.3M

#Preparing the data



```

import os, pathlib, shutil, random
from tensorflow import keras

#Setting up directories

batch_size = 32
base_dir = pathlib.Path("aclImdb")
val_dir = base_dir / "val"
train_dir = base_dir / "train"

for category in ("neg", "pos"):
    # Check if the directory already exists
    if not os.path.exists(val_dir / category):
        os.makedirs(val_dir / category)
    files = os.listdir(train_dir / category)
    random.Random(1337).shuffle(files)
    num_val_samples = int(0.2 * len(files))
    val_files = files[-num_val_samples:]
    for fname in val_files:
        # Check if the file already exists in the destination directory
        if not os.path.exists(val_dir / category / fname):
            shutil.move(train_dir / category / fname,
                        val_dir / category / fname)

#Creating Datasets from Directories

train_ds = keras.utils.text_dataset_from_directory(
    "aclImdb/train", batch_size=batch_size
)
val_ds = keras.utils.text_dataset_from_directory(
    "aclImdb/val", batch_size=batch_size
)
test_ds = keras.utils.text_dataset_from_directory(
    "aclImdb/test", batch_size=batch_size
)
text_only_train_ds = train_ds.map(lambda x, y: x)

🔍 Found 25000 files belonging to 2 classes.
   Found 5000 files belonging to 2 classes.
   Found 25000 files belonging to 2 classes.

#Preparing integer sequence datasets

from tensorflow.keras import layers

#Text Vectorization: Initialization and Adaptation

max_length = 600
max_tokens = 20000
text_vectorization = layers.TextVectorization(
    max_tokens=max_tokens,
    output_mode="int",
    output_sequence_length=max_length,
)
text_vectorization.adapt(text_only_train_ds)

#Integer Encoding Mapping for Datasets

int_train_ds = train_ds.map(
    lambda x, y: (text_vectorization(x), y),
    num_parallel_calls=4)
int_val_ds = val_ds.map(
    lambda x, y: (text_vectorization(x), y),
    num_parallel_calls=4)
int_test_ds = test_ds.map(
    lambda x, y: (text_vectorization(x), y),
    num_parallel_calls=4)

#A sequence model built on one-hot encoded vector sequences

import tensorflow as tf
from tensorflow import keras

```

```
from tensorflow.keras import layers
```

```
class OneHotEncodingLayer(layers.Layer):
    def __init__(self, depth, **kwargs):
        super().__init__(**kwargs)
        self.depth = depth

    def call(self, inputs):
        return tf.one_hot(inputs, depth=self.depth)
```

```
inputs = keras.Input(shape=(None,), dtype="int64")
# Wrap tf.one_hot in a custom layer
embedded = OneHotEncodingLayer(depth=max_tokens)(inputs)
x = layers.Bidirectional(layers.LSTM(32))(embedded)
x = layers.Dropout(0.5)(x)
outputs = layers.Dense(1, activation="sigmoid")(x)
model = keras.Model(inputs, outputs)
model.compile(optimizer="rmsprop",
              loss="binary_crossentropy",
              metrics=["accuracy"])
```

```
model.summary()
```

Model: "functional\_6"

Layer (type)	Output Shape	Param #
input_layer_6 (InputLayer)	(None, None)	0
one_hot_encoding_layer (OneHotEncodingLayer)	(None, None, 20000)	0
bidirectional (Bidirectional)	(None, 64)	5,128,448
dropout_5 (Dropout)	(None, 64)	0
dense_9 (Dense)	(None, 1)	65

Total params: 5,128,513 (19.56 MB)

```
from tensorflow import keras
from tensorflow.keras import layers
```

#Creating a simple neural network model for binary classification tasks to specify the maximum number of input tokens and the size of the hi

```
def get_model(max_tokens=20000, hidden_dim=16):
    inputs = keras.Input(shape=(max_tokens,))
    x = layers.Dense(hidden_dim, activation="relu")(inputs)
    x = layers.Dropout(0.5)(x)
    outputs = layers.Dense(1, activation="sigmoid")(x)
    model = keras.Model(inputs, outputs)
    model.compile(optimizer="rmsprop",
                  loss="binary_crossentropy",
                  metrics=["accuracy"])
    return model
```

```
# Define the model
model = get_model()
```

#Callback for Model Checkpointing

```
callbacks = [
    keras.callbacks.ModelCheckpoint("one_hot_bidir_lstm.keras",
                                    save_best_only=True)
]
```

```
import os
import random
import shutil
import pathlib
from tensorflow import keras
```

```

from tensorflow.keras import layers
import tensorflow as tf

# Define the base directory
base_dir = pathlib.Path("aclImdb")

# Define the directories for train, validation, and test sets
train_dir = base_dir / "train"
val_dir = base_dir / "val"
test_dir = base_dir / "test"

# Create directories for validation set
for category in ("neg", "pos"):
    os.makedirs(val_dir / category, exist_ok=True)
    files = os.listdir(train_dir / category)
    random.Random(1337).shuffle(files)
    num_val_samples = int(0.2 * len(files))
    val_files = files[-num_val_samples:]
    for fname in val_files:
        shutil.move(train_dir / category / fname, val_dir / category / fname)

#Data Preparation for Text Classification

batch_size = 32

train_ds = keras.utils.text_dataset_from_directory(
    train_dir, batch_size=batch_size, validation_split=0.2, subset="training", seed=1337
)
val_ds = keras.utils.text_dataset_from_directory(
    train_dir, batch_size=batch_size, validation_split=0.2, subset="validation", seed=1337
)
test_ds = keras.utils.text_dataset_from_directory(
    test_dir, batch_size=batch_size
)

🔗 Found 20000 files belonging to 2 classes.
Using 16000 files for training.
Found 20000 files belonging to 2 classes.
Using 4000 files for validation.
Found 25000 files belonging to 2 classes.

#Defining the TextVectorization layer
max_features = 20000
sequence_length = 600

vectorize_layer = layers.TextVectorization(
    max_tokens=max_features,
    output_mode='int',
    output_sequence_length=sequence_length,
)

#Adapting the TextVectorization layer on the training dataset
train_text = train_ds.map(lambda x, y: x)
vectorize_layer.adapt(train_text)

#Define the model
embedding_dim = 128
model = keras.Sequential([
    vectorize_layer,
    layers.Embedding(max_features + 1, embedding_dim),
    layers.Bidirectional(layers.LSTM(64)),
    layers.Dense(64, activation='relu'),
    layers.Dense(1)
])

#Compile the model
model.compile(
    loss=keras.losses.BinaryCrossentropy(from_logits=True),
    optimizer=keras.optimizers.Adam(1e-4),

```

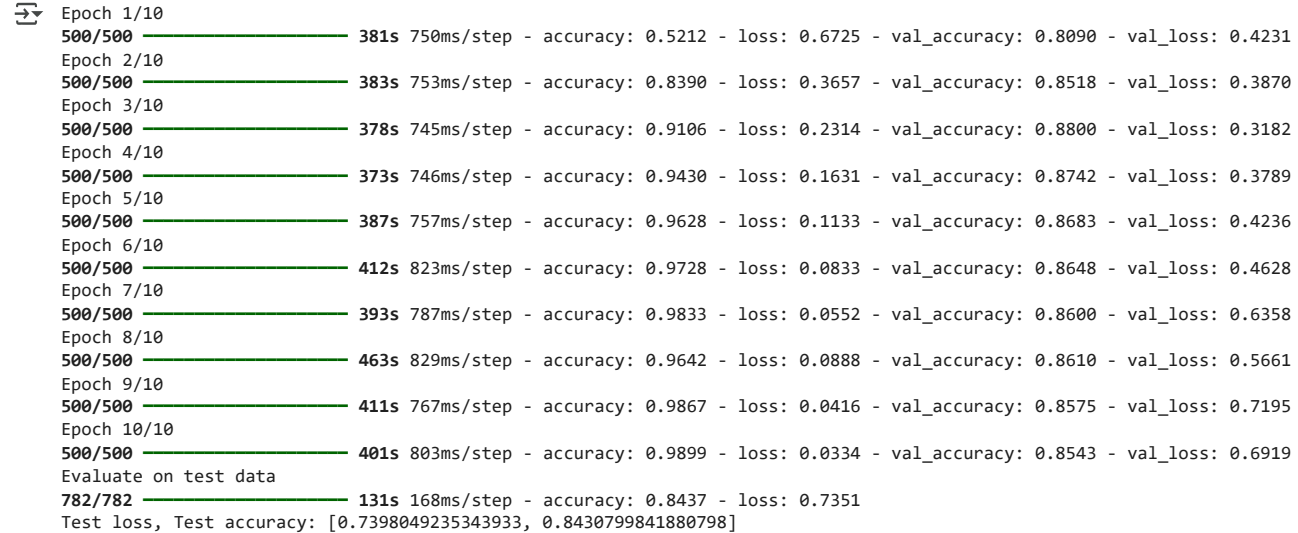
```

    metrics=['accuracy']
)

#Train the model
model.fit(train_ds, validation_data=val_ds, epochs=10)

#Evaluate the model
print("Evaluate on test data")
results = model.evaluate(test_ds)
print("Test loss, Test accuracy:", results)

```



Epoch	Time	Step	Accuracy	Loss	Val Accuracy	Val Loss
Epoch 1/10	381s	750ms/step	0.5212	0.6725	0.8090	0.4231
Epoch 2/10	383s	753ms/step	0.8390	0.3657	0.8518	0.3870
Epoch 3/10	378s	745ms/step	0.9106	0.2314	0.8800	0.3182
Epoch 4/10	373s	746ms/step	0.9430	0.1631	0.8742	0.3789
Epoch 5/10	387s	757ms/step	0.9628	0.1133	0.8683	0.4236
Epoch 6/10	412s	823ms/step	0.9728	0.0833	0.8648	0.4628
Epoch 7/10	393s	787ms/step	0.9833	0.0552	0.8600	0.6358
Epoch 8/10	463s	829ms/step	0.9642	0.0888	0.8610	0.5661
Epoch 9/10	411s	767ms/step	0.9867	0.0416	0.8575	0.7195
Epoch 10/10	401s	803ms/step	0.9899	0.0334	0.8543	0.6919
Evaluate on test data	782/782	131s 168ms/step	0.8437	0.7351		
Test loss, Test accuracy: [0.7398049235343933, 0.8430799841880798]						

```

max_tokens = 20000 # Defining the maximum number of tokens in your vocabulary
embedding_layer = layers.Embedding(input_dim=max_tokens, output_dim=256)

```

```

inputs = keras.Input(shape=(None,), dtype="int64")
embedded = layers.Embedding(input_dim=max_tokens, output_dim=256)(inputs)
x = layers.Bidirectional(layers.LSTM(32))(embedded)
x = layers.Dropout(0.5)(x)
outputs = layers.Dense(1, activation="sigmoid")(x)
model = keras.Model(inputs, outputs)
model.compile(optimizer="rmsprop",
              loss="binary_crossentropy",
              metrics=["accuracy"])

```

```

#Model definition with embedding layer for text classification

```

```

embedding_layer = layers.Embedding(input_dim=max_tokens, output_dim=256)

inputs = keras.Input(shape=(None,), dtype="int64")
embedded = layers.Embedding(input_dim=max_tokens, output_dim=256)(inputs)
x = layers.Bidirectional(layers.LSTM(32))(embedded)
x = layers.Dropout(0.5)(x)
outputs = layers.Dense(1, activation="sigmoid")(x)
model = keras.Model(inputs, outputs)
model.compile(optimizer="rmsprop",
              loss="binary_crossentropy",
              metrics=["accuracy"])

```

```

model.summary()

```

Model: "functional\_10"

Layer (type)	Output Shape	Param #
input_layer_10 ( <a href="#">InputLayer</a> )	(None, None)	0
embedding_5 ( <a href="#">Embedding</a> )	(None, None, 256)	5,120,000
bidirectional_3 ( <a href="#">Bidirectional</a> )	(None, 64)	73,984
dropout_8 ( <a href="#">Dropout</a> )	(None, 64)	0
dense_15 ( <a href="#">Dense</a> )	(None, 1)	65

Total params: 5,194,049 (19.81 MB)

Trainable params: 5,194,049 (19.81 MB)

#Directory paths for training, validation, and test data

```
train_dir = base_dir / "train"
val_dir = base_dir / "val"
test_dir = base_dir / "test"
```

#Model architecture definition with embedding layer, bidirectional LSTM, dropout, and dense layer

```
embedding_layer = layers.Embedding(input_dim=max_tokens, output_dim=256)
```

#Model input and layers configuration: embedding, bidirectional LSTM, dropout, and dense layers

```
inputs = keras.Input(shape=(None,), dtype="int64")
embedded = layers.Embedding(input_dim=max_tokens, output_dim=256)(inputs)
x = layers.Bidirectional(layers.LSTM(32))(embedded)
x = layers.Dropout(0.5)(x)
outputs = layers.Dense(1, activation="sigmoid")(x)
model = keras.Model(inputs, outputs)
```

#Model compilation with RMSprop optimizer, binary crossentropy loss, and accuracy metric

```
model.compile(optimizer="rmsprop",
              loss="binary_crossentropy",
              metrics=["accuracy"])
```

# Word Embedding Visualization  
from sklearn.decomposition import PCA


```
# Extract embeddings from the Embedding layer (index 1, not 0)
embedding_matrix = model.layers[1].get_weights()[0] # Changed from layers[0] to layers[1]
```

```
# Reduce dimensions using PCA
pca = PCA(n_components=2)
reduced_embeddings = pca.fit_transform(embedding_matrix)
```

```
# Plot the embeddings
plt.figure(figsize=(10, 10))
plt.scatter(reduced_embeddings[:, 0], reduced_embeddings[:, 1], alpha=0.7)
plt.title("Word Embedding Visualization (PCA)")
plt.xlabel("PCA Component 1")
plt.ylabel("PCA Component 2")
plt.show()
```



Word Embedding Visualization (PCA)



Could not connect to the reCAPTCHA service. Please check your internet connection and reload to get a reCAPTCHA challenge.